

Deep Generative Model


[original tutorial](#)

VAE


[HYL video](#)

Pixel RNN




R



G




B



PixelRNN


sequence of pixels

dependent on
previous pixels



Pixel CNN

PixelCNN




Kernel mask

1	1	1	1	1		
x_1	1	1	1	1	x_n	
1	1	x_i	0	0		
0	0	0	0	0		
0	0	0	0	0		
						x_{n^2}



WaveNet

WaveNet



$$p(\mathbf{x}) = \prod_{i=1}^{n^2} p(x_i | x_1, \dots, x_{i-1})$$



$$p(x|h) = \prod_{t=1}^T p(x_t | x_1, \dots, x_{t-1}, h)$$

$$p(\mathbf{x}) = \prod_{t=1}^T p(x_t | x_1, \dots, x_{t-1})$$

直接用聲音的原始波形來進行訓練，生成聲波訊號。

sigmoid 作為activation function

另外可以加入h作為conditional WaveNet



Challenges for Graph Generation

- The structures and sizes of graphs are different
- No orders between the nodes
- Discrete


Graph VAE

[paper: arxiv 1802.03480](#)

9 Feb 2018


Graph VAE

GraphVAE: Towards Generation of Small Graphs Using Variational Autoencoders



Graph VAE


GraphVAE: Towards Generation of Small Graphs Using Variational Autoencoders



Graph VAE

$$\mathcal{L}(\phi, \theta; G) = \underbrace{\mathbb{E}_{q_\phi(\mathbf{z}|G)}[-\log p_\theta(G|\mathbf{z})]}_{\text{reconstruction loss}} + \underbrace{\text{KL}[q_\phi(\mathbf{z}|G)||p(\mathbf{z})]}_{\text{regularization}} \quad (1)$$

GraphVAE: Towards Generation of Small Graphs Using Variational Autoencoders



Graph VAE

\tilde{A} : predicted adjacency [node prob & edge prob]

sigmoid


\tilde{E} : edge attribute, (class prob for edges)

softmax

\tilde{F} : node attribute, (class prob for nodes)

softmax

GraphVAE: Towards Generation of Small Graphs Using Variational Autoencoders




Graph VAE

$$\begin{aligned}
 S((i, j), (a, b)) = & \\
 = & (E_{i,j}^T, \tilde{E}_{a,b}, \cdot) A_{i,j} \tilde{A}_{a,b} \tilde{A}_{a,a} \tilde{A}_{b,b} [i \neq j \wedge a \neq b] + \\
 + & (F_{i,\cdot}^T, \tilde{F}_{a,\cdot}) \tilde{A}_{a,a} [i = j \wedge a = b]
 \end{aligned} \quad (4)$$

similarity
function

$$X \in \{0, 1\}^{k \times n}$$

GraphVAE: Towards Generation of Small Graphs Using Variational Autoencoders



Graph VAE

$$\mathcal{L}(\phi, \theta; G) = \underbrace{\mathbb{E}_{q_\phi(\mathbf{z}|G)}[-\log p_\theta(G|\mathbf{z})]}_{\text{reconstruction loss}} + \underbrace{\text{KL}[q_\phi(\mathbf{z}|G)||p(\mathbf{z})]}_{\text{regularization}} \quad (1)$$

$$X \in \{0, 1\}^{k \times n}$$

input adjacency matrix is mapped to the predicted graph

$$\bar{A}' = X A X^T$$

predicted **node** attribute matrix and slices of **edge** attribute matrix are transferred to the input graph as

$$\tilde{F}' = X^T \tilde{F}, \quad \tilde{E}'_{\cdot, \cdot, l} = X^T \tilde{E}_{\cdot, \cdot, l} X.$$

The overall **reconstruction loss** is a weighed sum of the previous terms:

$$\begin{aligned} -\log p(G|\mathbf{z}) = & -\lambda_A \log p(A'|\mathbf{z}) - \lambda_F \log p(F|\mathbf{z}) - \\ & - \lambda_E \log p(E|\mathbf{z}) \end{aligned} \quad (3)$$

Graph VAE

$$X \in \{0, 1\}^{k \times n}$$

obtain from graph matching, a binary assignment matrix

input adjacency matrix is mapped to the predicted graph $\bar{A}' = \bar{X} A X^T$

predicted **node** attribute matrix and slices of **edge** attribute matrix are transferred to the input graph as

$$\tilde{F}' = X^T \tilde{F}$$

$$\tilde{E}'_{:,l} = X^T \tilde{E}_{:,l} X$$

$$\mathcal{L}(\phi, \theta; G) =$$

$$= \mathbb{E}_{q_\phi(\mathbf{z}|G)} [-\log p_\theta(G|\mathbf{z})] + \text{KL}[q_\phi(\mathbf{z}|G) || p(\mathbf{z})] \quad (1)$$

reconstruction loss

regularization

The **maximum likelihood** estimates, i.e. cross-entropy, of respective variables are as follows:

$$\log p(A'|\mathbf{z}) =$$

$$= 1/k \sum_a A'_{a,a} \log \tilde{A}_{a,a} + (1 - A'_{a,a}) \log(1 - \tilde{A}_{a,a}) +$$

$$+ 1/k(k-1) \sum_{a \neq b} A'_{a,b} \log \tilde{A}_{a,b} + (1 - A'_{a,b}) \log(1 - \tilde{A}_{a,b})$$

$$\log p(F|\mathbf{z}) = 1/n \sum_i \log F_{i,\cdot}^T \tilde{F}'_{i,\cdot}$$

$$\log p(E|\mathbf{z}) = 1/(||A||_1 - n) \sum_{i \neq j} \log E_{i,j,\cdot}^T \tilde{E}'_{i,j,\cdot}, \quad (2)$$

The overall **reconstruction loss** is a weighed sum of the previous


$$-\log p(G|\mathbf{z}) = -\lambda_A \log p(A'|\mathbf{z}) - \lambda_F \log p(F|\mathbf{z}) - \lambda_E \log p(E|\mathbf{z}) \quad (3)$$

JTVAE

[paper: arxiv 1802.04364](#)


29 Mar 2019

JTVAE - Junction Tree



A tree decomposition maps a graph G into a junction tree by contracting certain vertices into a single node so that G becomes cycle-free.

JTVAE - Graph Encoder



$$\nu_{uv}^{(l)} = \tau(\mathbf{W}_1^g \mathbf{x}_u + \mathbf{W}_2^g \mathbf{x}_{uv} + \mathbf{W}_3^g \sum_{w \in N(u) \setminus v} \nu_{wu}^{(l-1)}) \quad (1)$$


$$\mathbf{h}_u = \tau(\mathbf{U}_1^g \mathbf{x}_u + \sum_{v \in N(u)} \mathbf{U}_2^g \nu_{vu}^{(T)}) \quad (2)$$

The final graph representation is $\mathbf{h}_G = \sum_i \mathbf{h}_i / |V|$.

message passing through edges

compute the messages of latent vector of each node

JTVAE - Tree Encoder



top-down, bottom up
GRU over tree nodes


$$m_{ij} = \text{GRU}(x_i, \{m_{ki}\}_{k \in N(i) \setminus j}) \quad (3)$$

After the message passing, we obtain the latent representation of each node h_i by aggregating its inward messages:

$$h_i = \tau(\mathbf{W}^o \mathbf{x}_i + \sum_{k \in N(i)} \mathbf{U}^o m_{ki}) \quad (9)$$

The final tree representation is $h_{\mathcal{T}_G} = h_{root}$,

JTVAE - Tree Decoder



topological prediction:

$$p_t = \sigma(\mathbf{u}^d \cdot \tau(\mathbf{W}_1^d \mathbf{x}_{i_t} + \mathbf{W}_2^d \mathbf{z}_\mathcal{T} + \mathbf{W}_3^d \sum_{(k, i_t) \in \tilde{\mathcal{E}}_t} \mathbf{h}_{k, i_t})) \quad (11)$$


new child's label prediction:

$$\mathbf{q}_j = \text{softmax}(\mathbf{U}^l_\tau(\mathbf{W}_1^l \mathbf{z}_\mathcal{T} + \mathbf{W}_2^l \mathbf{h}_{ij})) \quad (12)$$

Let p and q be the ground truth topological and label values, the decoder minimizes the following cross entropy loss:

$$\mathcal{L}_c(\mathcal{T}) = \sum_t \mathcal{L}^d(p_t, \hat{p}_t) + \sum_j \mathcal{L}^l(\mathbf{q}_j, \hat{\mathbf{q}}_j) \quad (13)$$

JTVAE - Graph Decoder



Let $\mathcal{G}(\mathcal{T})$ be the set of graphs whose junction tree is \mathcal{T} . Decoding graph \hat{G} from $\hat{\mathcal{T}} = (\hat{\mathcal{V}}, \hat{\mathcal{E}})$ is a structured prediction:

$$\hat{G} = \arg \max_{G' \in \mathcal{G}(\hat{\mathcal{T}})} \boxed{f^a(G')} \quad (14)$$

scoring function

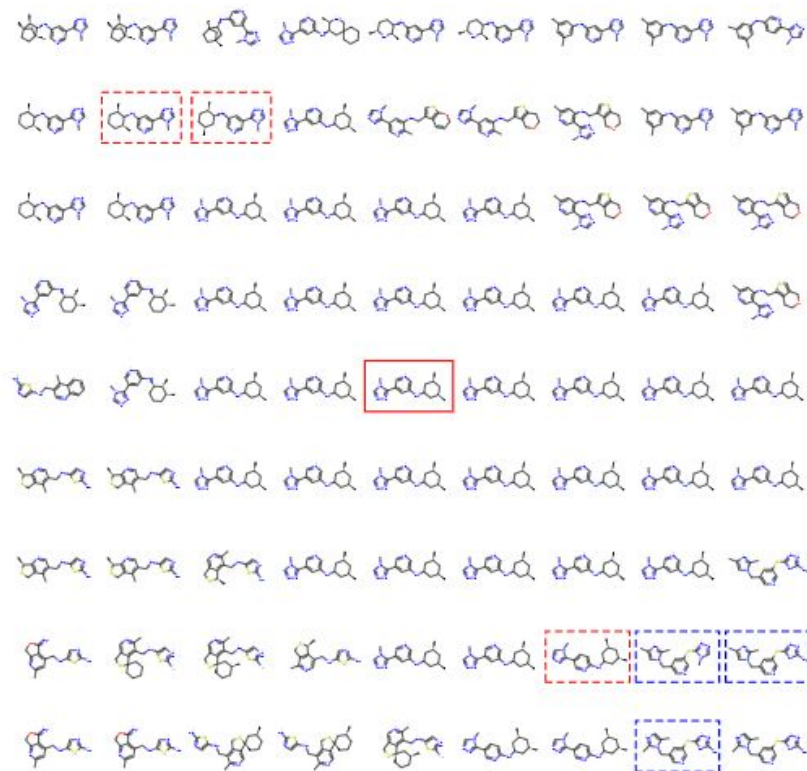
$$f_i^a(G_i) = \mathbf{h}_{G_i} \cdot \mathbf{z}_G$$

$$\mathcal{L}_g(G) = \sum_i \left[f^a(G_i) - \log \sum_{G'_i \in \mathcal{G}_i} \exp(f^a(G'_i)) \right] \quad (16)$$

maximize the log-likelihood of predicting correct subgraphs G_i of the ground true graph G at each tree node. We again apply teacher forcing, i.e. we feed the graph decoder with ground truth trees as input.

JTVAE - Graph Decoder


same tree structure
but different moleculars



MolGAN


[paper: arxiv 1805.11973](https://arxiv.org/abs/1805.11973)

30 May 2018



MolGAN

MolGAN: An implicit generative model for small molecular graphs



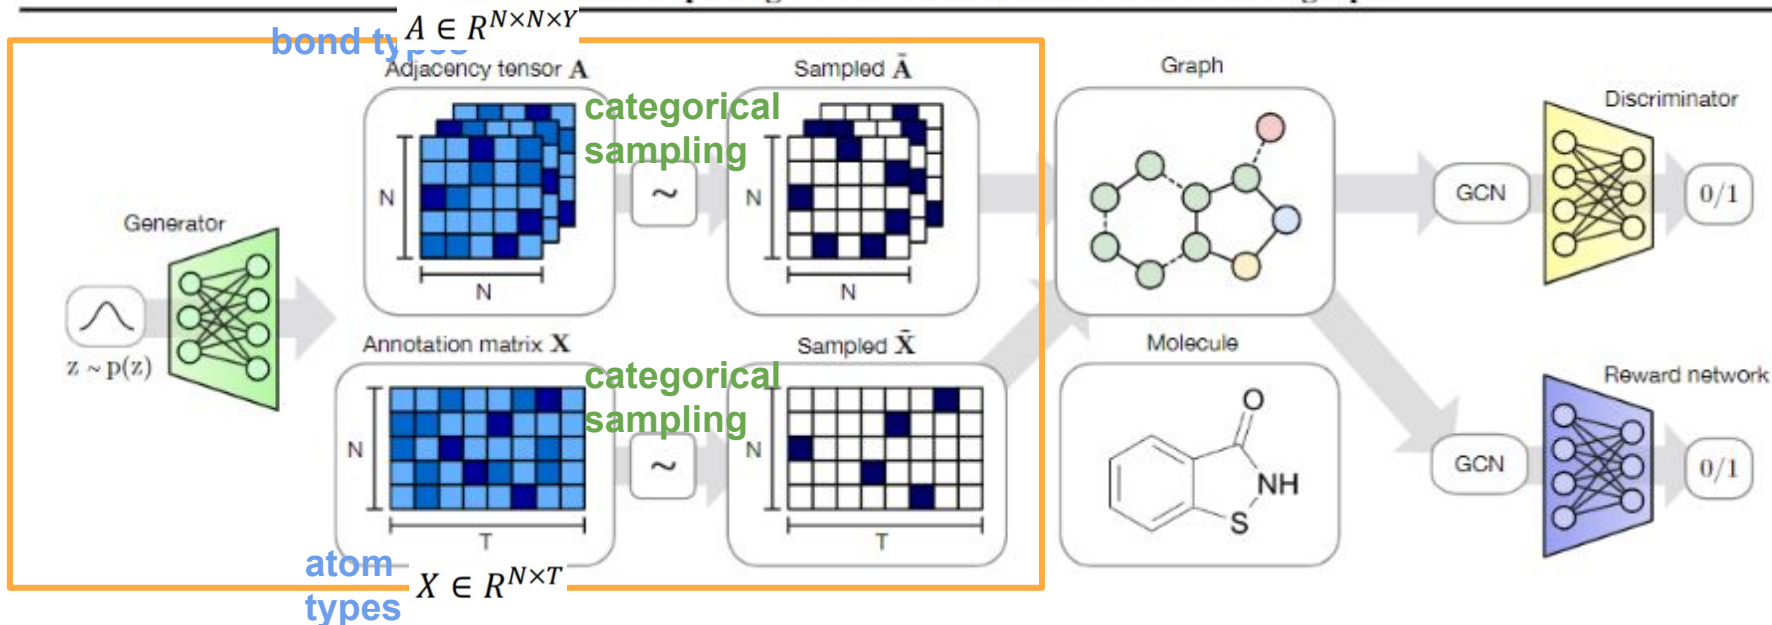
implicit, likelihood-free.
using RL for desired properties.

loss

$$L(\theta) = \lambda \cdot L_{WGAN} + (1 - \lambda) \cdot L_{RL}$$

MolGAN - Generator


MolGAN: An implicit generative model for small molecular graphs



Generator: predicts the entire graph at once using a simple multi-layer perceptron (MLP). While this limits our study to graphs of a pre-chosen maximum size, we find that it is significantly faster and easier to optimize.

MolGAN - Discriminator & Reward Network

MolGAN: An implicit generative model for small molecular graphs

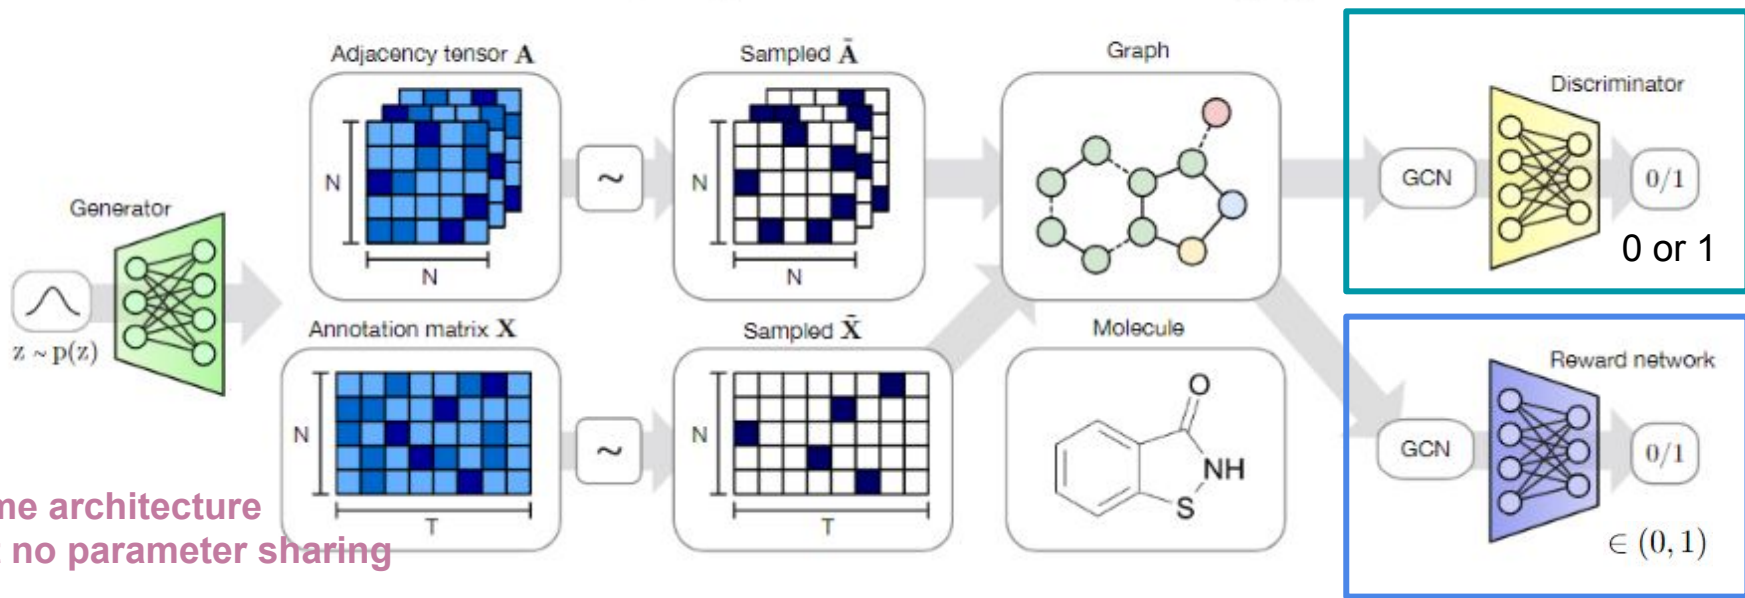


Relational - GCN
support for
multiple edge types

MolGAN - Discriminator & Reward Network

MolGAN: An implicit generative model for small molecular graphs

input: graph
output: scalar value



Discriminator: samples from the dataset and the generator \rightarrow learns to distinguish them

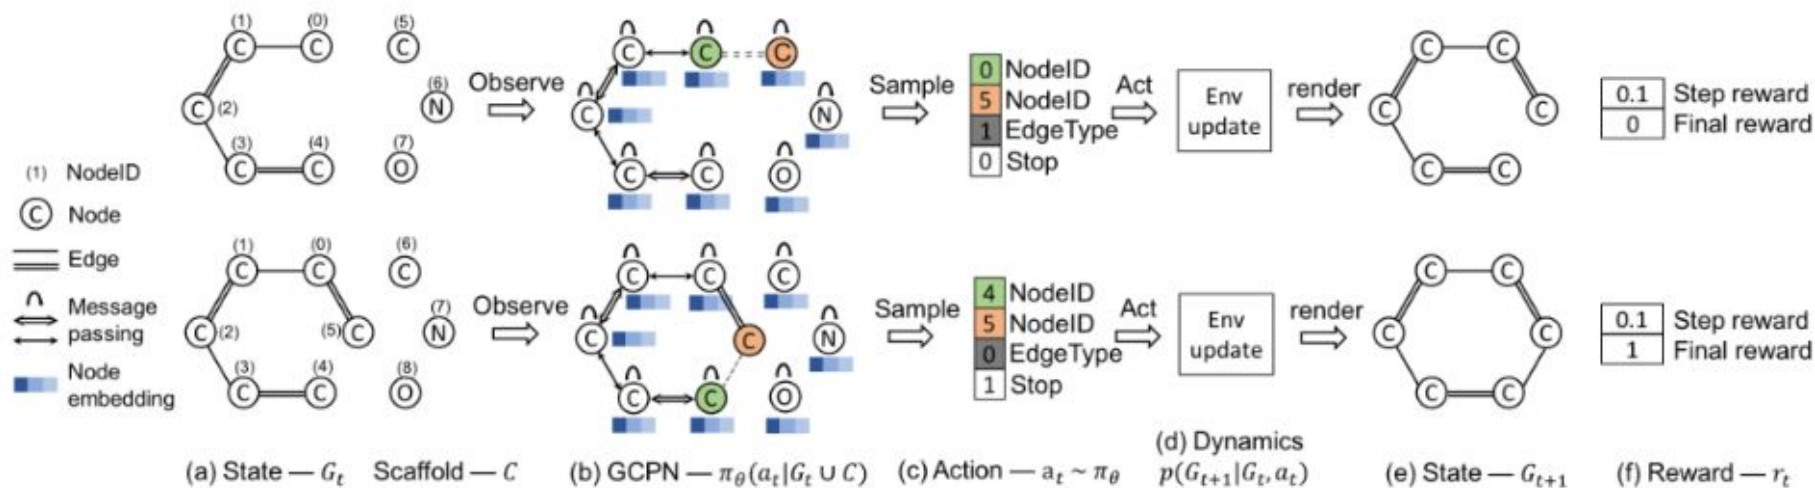
Reward Network: dataset and generated samples are inputs \rightarrow assigns scores to them (using [RDKit](#))

GCPN

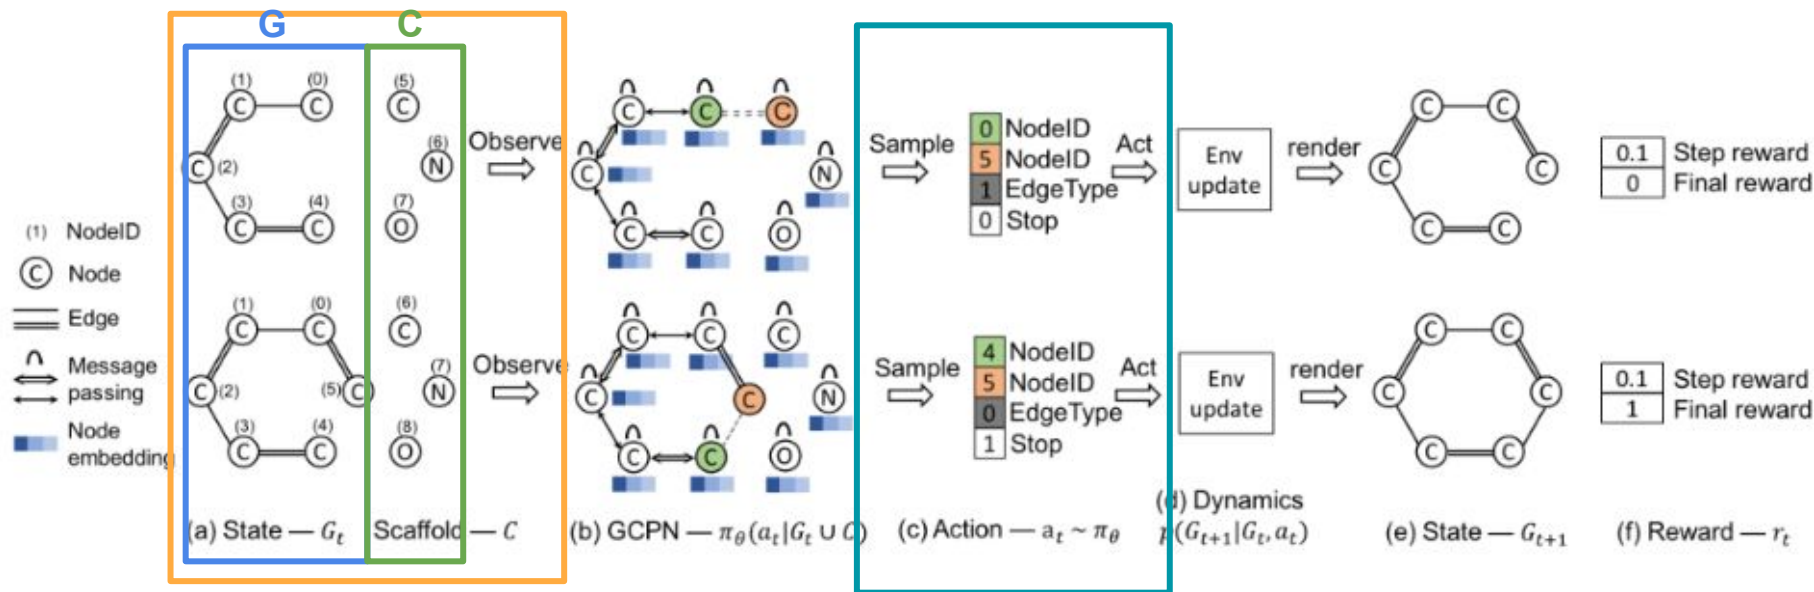
paper

2018

GCPN



GCPN - State Space & Action Space

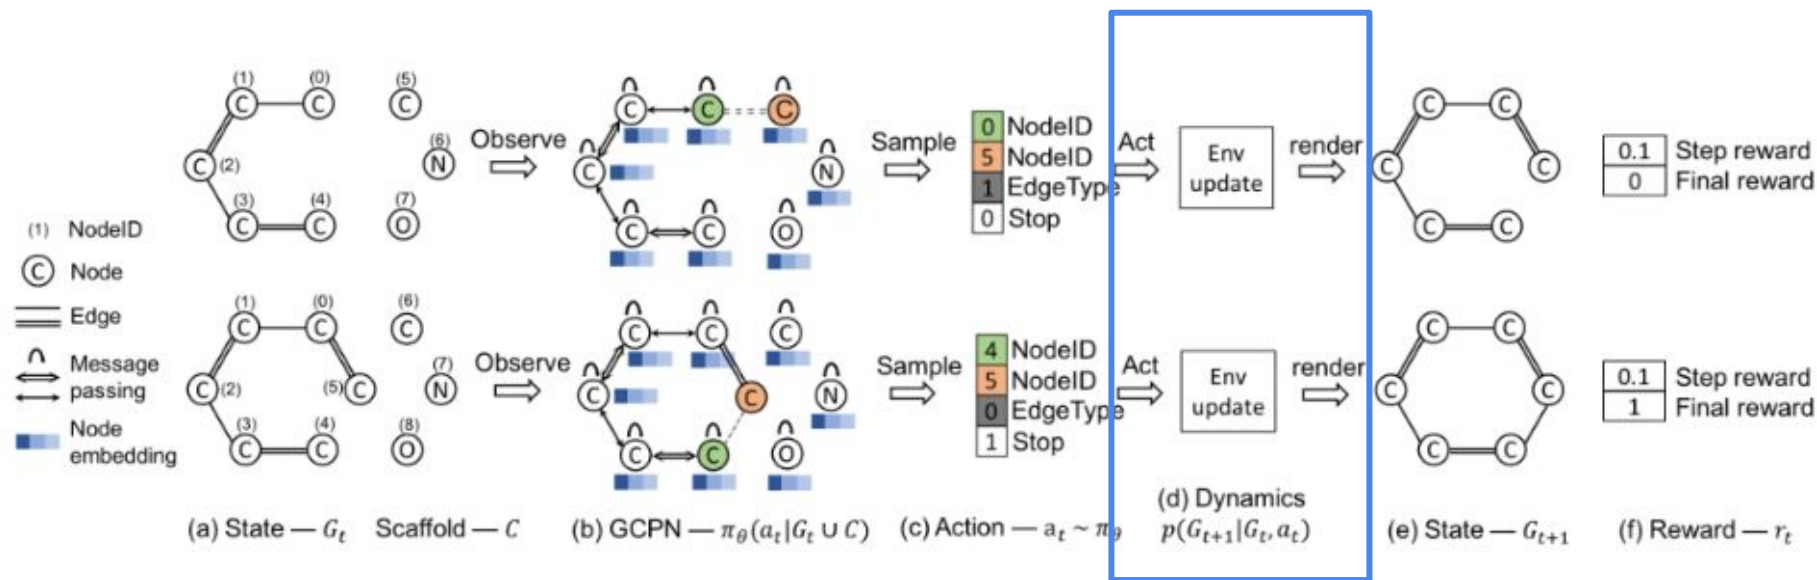


State space: 初始值給定一個碳元素 node

Action space: 視作link prediction a set of scaffold subgraphs $\{C_1, \dots, C_s\}$

- Connect a new subgraph C_i to a node in G_t or
- connecting existing nodes within graph G_t


GCPN - State Transition Dynamics



State Transition Dynamics: 考量Domain-specific rules, step-wise valency check in incompleted graph

- obey the given rules or
- be rejected and the state remains unchanged

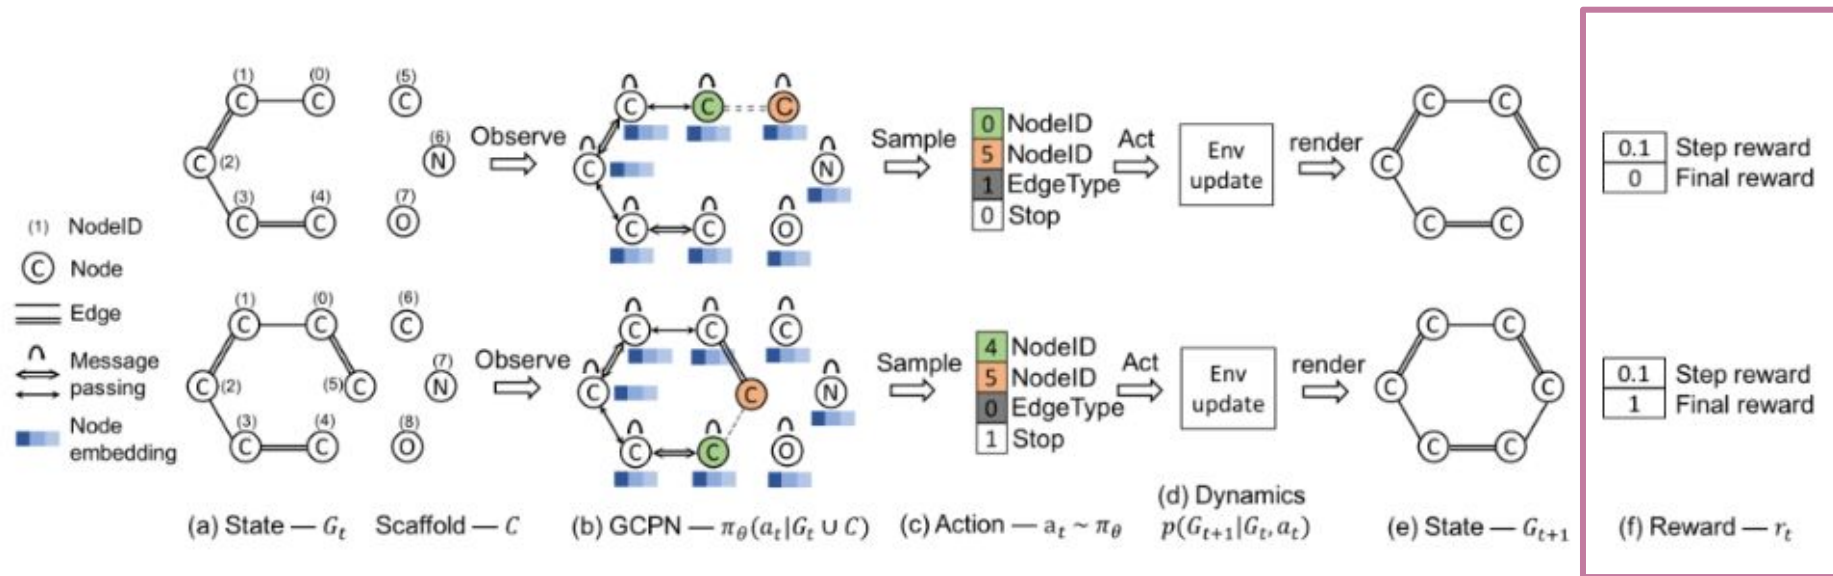
GCPN - Reward Design



Reward Design:

- intermediate rewards:
 - step-wise validity rewards and adversarial rewards
- final rewards
 - a sum over domain-specific rewards and adversarial rewards

GCPN - Reward Design



Reward Design: the adversarial rewards

$$\min_{\theta} \max_{\phi} V(\pi_{\theta}, D_{\phi}) = \mathbb{E}_{x \sim p_{data}} [\log D_{\phi}(x)] + \mathbb{E}_{x \sim \pi_{\theta}} [\log D_{\phi}(1 - x)]$$

GCPN - Graph Convolutional Policy Network


computing node embeddings using GCN

then make an action prediction

$$a_t = \text{CONCAT}(a_{\text{first}}, a_{\text{second}}, a_{\text{edge}}, a_{\text{stop}}) \quad (3)$$

$$\begin{aligned} f_{\text{first}}(s_t) &= \text{SOFTMAX}(m_f(X)), & a_{\text{first}} &\sim f_{\text{first}}(s_t) \in \{0, 1\}^n \\ f_{\text{second}}(s_t) &= \text{SOFTMAX}(m_s(X_{a_{\text{first}}}, X)), & a_{\text{second}} &\sim f_{\text{second}}(s_t) \in \{0, 1\}^{n+c} \\ f_{\text{edge}}(s_t) &= \text{SOFTMAX}(m_e(X_{a_{\text{first}}}, X_{a_{\text{second}}}), & a_{\text{edge}} &\sim f_{\text{edge}}(s_t) \in \{0, 1\}^b \\ f_{\text{stop}}(s_t) &= \text{SOFTMAX}(m_t(\text{AGG}(X))), & a_{\text{stop}} &\sim f_{\text{stop}}(s_t) \in \{0, 1\} \end{aligned} \quad (4)$$

GCPN - Policy Gradient Training



Proximal Policy Optimization (PPO) & objective function during pretraining (supervised)

$$\max L^{\text{CLIP}}(\theta) = \mathbb{E}_t[\min(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)\hat{A}_t)], r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)} \quad (5)$$

$\min L^{\text{EXPERT}}(\theta) = -\log(\pi_\theta(a_t|s_t))$ where (st;at) pairs are obtained from ground truth molecules.

GCPN - Experiments

specified properties are optimized

Property Optimization. The task is to generate novel molecules whose specified molecular properties are optimized. This can be useful in many applications such as drug discovery and materials science, where the goal is to identify molecules with highly optimized properties of interest.

properties are as close to the target score

Property Targeting. The task is to generate novel molecules whose specified molecular properties are as close to the target scores as possible. This is crucial in generating virtual libraries of molecules with properties that are generally suitable for a desired application. For example, a virtual molecule library for drug discovery should have high drug-likeness and synthesizability.

Constrained Property Optimization. The task is to generate novel molecules whose specified molecular properties are optimized, while also containing a specified molecular substructure. This can be useful in lead optimization problems in drug discovery and materials science, where we want to make modifications to a promising lead molecule and improve its properties [2].

**specified properties are optimized,
and also containing specified substructure**

GCPN[98]利用RL生成目標導向的分子圖，以處理不可導目標和約束。實驗結果證明了GCPN在各種圖生成問題中的有效性。

MoIGAN[99]也採用了類似的思想，即使用RL生成分子圖。MoIGAN建議直接生成完整的圖，而不是通過一系列的動作來生成圖，這對小分子很有效。